

Fast Multidimensional Signal Processing with Shearlab.jl

Héctor Andrade Loarca


TU Berlin, BMS



Fourier Transform.

```
In [1]: using PyPlot
```

```
In [2]: # Create and plot signal
t=0:2/256:2;
f=sin(2*pi*t)+2*(cos(6*pi*t))*2;
plot(t,f)
xlabel("t")
ylabel("f(t)")
```




```
Out[2]: PyObject <matplotlib.text.Text object at 0x31ff4c550>
```

```
In [3]: # Compute fast fourier transform
@time F = fft(f);
```

```
0.511862 seconds (414.51 k allocations: 17.621 MB, 1.52% gc time)
```


```
In [4]: plot(t,F)
        xlabel("ω")
        ylabel("F(ω)")
```



```
/Users/hector/.julia/v0.5/Conda/deps/usr/lib/python2.7/site-packages/numpy/core/numeric.py:531: ComplexWarning: Casting complex values to real discards the imaginary part
  return array(a, dtype, copy=False, order=order)
```

```
Out[4]: PyObject <matplotlib.text.Text object at 0x320064d90>
```

```
In [5]: plot(t,ifft(F))
        xlabel("t")
        ylabel("ifft(F)(t)")
```



Out[5]: PyObject <matplotlib.text.Text object at 0x31feaa990>

Wavelet Transform

```
In [6]: # using Shearlab
        # Load the Pkg
        push!(LOAD_PATH,pwd()*"/../..../Shearlab.jl/src")
        import Shearlab
        reload("Shearlab")
```

WARNING: replacing module Shearlab

The 2D wavelet transform

$$d_j^k[n] = \langle f, \psi_{j,n}^k \rangle$$

for scales $j \in \mathbb{Z}$, position $n \in \mathbb{Z}^2$ and orientation $k \in \{H, V, D\}$.

And wavelet atoms defined by scaling and translating three mother atoms $\{\psi^H, \psi^V, \psi^D\}$:

$$\psi_{j,n}^k(x) = \frac{1}{2^j} \psi^k\left(\frac{x - 2^j n}{2^j}\right)$$

Defined by tensor product of a 1-D wavelet function $\psi(t)$ and a 1-D scaling function $\phi(t)$

$$\psi^H(x) = \phi(x_1)\psi(x_2), \psi^V(x) = \psi(x_1)\phi(x_2) \text{ and } \psi^D(x) = \psi(x_1)\psi(x_2).$$

Translated into high-pass and low-pass filters

$$g[n] = \frac{1}{\sqrt{2}} \langle \psi(t/2), \phi(t-n) \rangle \text{ and } h[n] = \frac{1}{\sqrt{2}} \langle \phi(t/2), \phi(t-n) \rangle.$$

1D

Compute the low pass signal $a \in \mathbb{R}^{N/2}$ and the high pass signal $d \in \mathbb{R}^{N/2}$ as

$$a = (f * h) \downarrow 2 \text{ and } d = (f * g) \downarrow 2$$

where the sub-sampling is defined as

$$(u \downarrow 2)[k] = u[2k].$$

where $g[n] = (-1)^{1-n}h[1-n]$.

When the filters are orthogonal transform the inverse will be

$$(a \uparrow 2) * \tilde{h} + (d \uparrow 2) * \tilde{g} = f$$

where $\tilde{h}[n] = h[-n]$ (computed modulo N) and $(u \uparrow 2)[2n] = u[n]$ and $(u \uparrow 2)[2n+1] = 0$.

2D

One needs to perform filtering/downsampling in the different directions

$$\tilde{a}_{j-1} = (a_j *^H h) \downarrow^{2,H} \text{ and } \tilde{d}_{j-1} = (a_j *^H g) \downarrow^{2,H}.$$

Here, the operator $*^H$ and $\downarrow^{2,H}$ are defined by applying $*$ and \downarrow^2 to each column of the matrix.

$$a_{j-1} = (\tilde{a}_j *^V h) \downarrow^{2,V} \text{ and } d_{j-1}^V = (\tilde{a}_j *^V g) \downarrow^{2,V},$$

$$d_{j-1}^H = (\tilde{d}_j *^V h) \downarrow^{2,V} \text{ and } d_{j-1}^D = (\tilde{d}_j *^V g) \downarrow^{2,V}.$$

```
In [7]: using Wavelets
```

```
In [8]: # Pick Daubechie filter and its mirror
h = Shearlab.filt_gen(WT.db2)
g = Shearlab.mirror(h)
```

```
Out[8]: 5-element Array{Float64,1}:
-0.0
-0.12941
-0.224144
 0.836516
-0.482963
```

```
In [9]: # Read Data
n = 1024;
# The path of the image
name = "../.../ShearLab.jl/data_samples/lena.jpg";
f = Shearlab.load_image(name, n,n,1);
```

```
In [10]: #Rescale image in [0,1] summing the 3 arrays in the RGB format  
         f = Shearlab.rescale(sum(f,3));  
         # Reduce one dimension  
         f = f[:, :, 1];  
         Shearlab.imageplot(f);
```




```


In [11]: # Different scales
Jmax = round(Int64,log2(n))-1;
Jmin = 1;
fW = copy(f);
clf;
figure(figsize=(10,10));
for j=Jmax:-1:Jmin
    A = fW[1:2^(j+1),1:2^(j+1)];
    for d=1:2
        Coarse = Shearlab.subsampling(Shearlab.cconvol(A,h,d),d);
        Detail = Shearlab.subsampling(Shearlab.cconvol(A,g,d),d);
        A = cat(d, Coarse, Detail );
    end
    fW[1:2^(j+1),1:2^(j+1)] = A;
    j1 = Jmax-j;
    if j1<4
        Shearlab.imageplot(A[1:2^j,2^j+1:2^(j+1)], "Horizontal, j=$j", 3,4,
        Shearlab.imageplot(A[2^j+1:2^(j+1),1:2^j], "Vertical, j=$j", 3,4, j1
        Shearlab.imageplot(A[2^j+1:2^(j+1),2^j+1:2^(j+1)], "Diagonal, j=$j",
    end
end

```


Horizontal, j=9



Horizontal, j=8




Horizontal, j=7




Horizontal, j=6




Vertical, j=9



Vertical, j=8




Vertical, j=7




Vertical, j=6




Diagonal, j=9




Diagonal, j=8



Diagonal, j=7



Diagonal, j=6



```

In [12]: f1 = copy(fW);
         clf;
         figure(figsize=(10,10));
         for j=Jmin:Jmax
             A = f1[1:2^(j+1),1:2^(j+1)];
             for d=1:2
                 if d==1
                     Coarse = A[1:2^j,:];
                     Detail = A[2^j+1:2^(j+1),:];
                 else
                     Coarse = A[:,1:2^j];
                     Detail = A[:,2^j+1:2^(j+1)];
                 end
                 Coarse = Shearlab.cconvol(Shearlab.upsampling(Coarse,d),Shearlab.rev
                 Detail = Shearlab.cconvol(Shearlab.upsampling(Detail,d),Shearlab.rev
                 A = Coarse + Detail;
                 j1 = Jmax-j;
                 if j1>0 && j1<5
                     Shearlab.imageplot(A, "Partial reconstruction, j=$j", 2,2,j1);
                 end
             end
             f1[1:2^(j+1),1:2^(j+1)] = A;
         end

```


Partial reconstruction, $j=8$ Partial reconstruction, $j=7$ Partial reconstruction, $j=6$ Partial reconstruction, $j=5$ 

```

In [13]: # number of kept coefficients
m = round(n^2/16);
# compute the threshold T
Jmin = 1;
fW = Shearlab.perform_wavortho_transf(f,Jmin,+1, h);
# select threshold
v = sort(abs(fW[:]));
if v[1]<v[n^2]
    v = Shearlab.reverse(v);
end
# inverse transform
T = v[Int(m)];
fWT = fW .* (abs(fW).>T);2
# inverse
fnlin = Shearlab.perform_wavortho_transf(fWT,Jmin,-1, h);
# display
clf;
figure(figsize=(10,10));
u1 = @sprintf("Original");
u2 = @sprintf("Thresholding, SNR=%3fdB", Shearlab.snr(f,fnlin));
Shearlab.imageplot(Shearlab.clamp(f),u1, 1,2,1 );
Shearlab.imageplot(Shearlab.clamp(fnlin),u2, 1,2,2 );

```

Original



Thresholding, SNR=36.322dB



Non-separable Shearlet Transform

$$\text{DST}_{j,k,m}^{2D}(f) = \left(\overline{\psi_{j,k}^d} * f_J \right) \left(2^j A_{2^j}^{-1} M_{c_j} m \right) \text{ for } j = 0, 1, \dots, J-1$$

with generating function

$$\hat{\psi}(\xi) = P(\xi_1/2, \xi_2/2) \widehat{\psi_1 \otimes \phi_1}(\xi)$$

where

$$\psi_{j,k}^d = S_{k/2^{j/2}}^d(p_j * W_j)$$

A_{2^j} is the parabolic scaling matrix given by

$$A_{2^j} = \begin{pmatrix} 2^j & 0 \\ 0 & 2^{j/2} \end{pmatrix}$$

and the Shearing transform is given by

$$S_k = \begin{pmatrix} 1 & k \\ 0 & 1 \end{pmatrix}$$

The scaling and wavelet filters now in each scale are:

$$\hat{h}_j(\xi_1) = \prod_{k=0}^{j-1} \hat{h}(2^k \xi_1), \hat{g}_j(\xi_1) = \hat{g}\left(\frac{2^j \xi_1}{2}\right) \hat{h}_{j-1}(\xi_1)$$

and the directional filter transform

$$p_j(\xi_1, \xi_2) = P(\widehat{2^{J-j-1}\xi_1}, 2^{J-j/2}\xi_2) \implies p_j * W_j = p_j * g_{J-j} \otimes h_{J-j}$$

The digital shearing to maintain the domain grid will be

$$S_{k/2^{j/2}}(x) = \left((x_{\uparrow 2^{j/2}} *_1 h_{j/2}) (S_k \cdot) *_1 \overline{h_{j/2}} \right)_{\downarrow 2^{j/2}}$$

for $j \in \{0, J-1\}$ and $|k| \leq \lceil 2^{j/2} \rceil$.

```
In [14]: # Read Data
n = 512;
# The path of the image
name = "../../../ShearLab.jl/data_samples/lena.jpg";
data_nopar = Shearlab.load_image(name, n);
data_par = Shearlab.load_image(name, n,n,1);
```

```
In [15]: # Reduce one dimension
data_nopar = data_nopar[:, :, 1];
data_par = data_par[:, :, 1]
Shearlab.imageplot(data_nopar);
```



```
In [16]: Shearlab.imageplot(data_par);
```



```
In [17]: # Size of the images  
sizeX_nopar = size(data_nopar,1);  
sizeY_nopar = size(data_nopar,2);  
sizeX_par = size(data_par,1);  
sizeY_par = size(data_par,2);
```

```
In [18]: # Set the variables for the Shearlet transform  
rows_nopar = sizeX_nopar;  
cols_nopar = sizeY_nopar;  
rows_par = sizeX_par;  
cols_par = sizeY_par;  
X_nopar = data_nopar;  
X_par = data_par;
```

```
In [19]: # No. of scales  
nScales = 4;  
shearLevels = ceil((1:nScales)/2)  
scalingFilter = Shearlab.filt_gen("scaling_shearlet");  
directionalFilter = Shearlab.filt_gen("directional_shearlet");  
waveletFilter = Shearlab.mirror(scalingFilter);  
scalingFilter2 = scalingFilter;  
full = 0;
```

Generation of Shearlet System

```
In [24]: # Compute the corresponding shearlet system without gpu  
@time shearletSystem_nopar= Shearlab.getshearletsystem2D(rows_nopar,cols_nopar,  
                                                         directionalFilter,  
                                                         scalingFilter,0);
```


8.314009 seconds (154.89 k allocations: 5.024 GB, 45.95% gc time)

```
In [21]: using ArrayFire
```

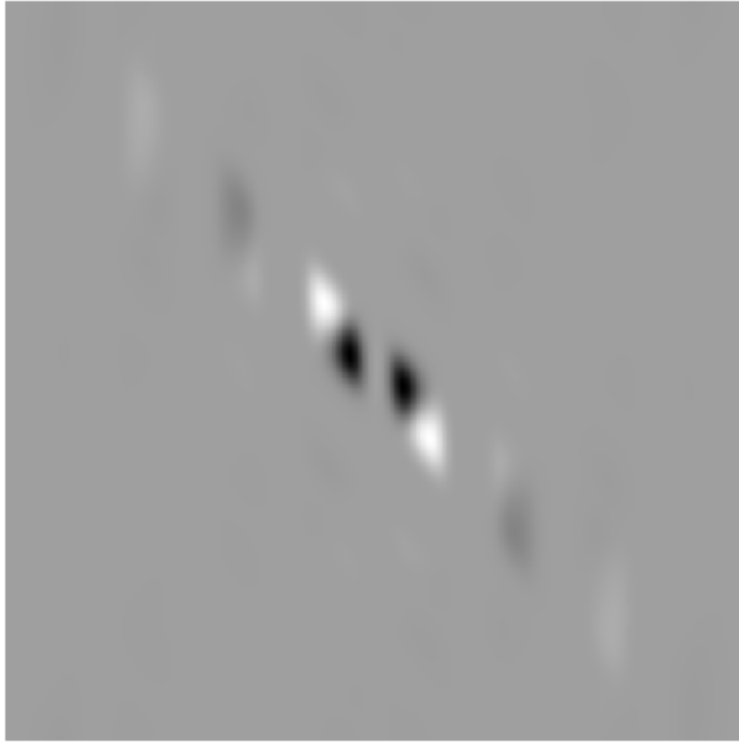
```
In [25]: # Compute the corresponding shearlet system without gpu  
@time shearletSystem_par = Shearlab.getshearletsystem2D(rows_par,cols_par,ns  
                                                         directionalFilter,  
                                                         scalingFilter,1);
```

2.425434 seconds (173.99 k allocations: 506.944 MB, 15.23% gc time)


```
In [26]: shearlet1 = Array(shearletSystem_nopar.shearlets[:, :, 1]);  
Shearlab.imageplot(real(shearlet1));
```



```
In [27]: # Comparison with the generated by arrayfire
shearlet1 = Array(shearletSystem_par.shearlets[:, :, 1]);
Shearlab.imageplot(real(shearlet1));
```



```
In [28]: shearlet3 = Array(shearletSystem_nopar.shearlets[:, :, 3]);  
Shearlab.imageplot(real(shearlet3))
```



```
In [29]: shearlet10 = Array(shearletSystem_nopar.shearlets[:, :, 10]);  
Shearlab.imageplot(real(shearlet10))
```



```
In [30]: shearlet17 = Array(shearletSystem_nopar.shearlets[:, :, 49]);  
        Shearlab.imageplot(real(shearlet17))
```




Computation of shearlet coefficients

```
In [33]: # Compute the coefficients  
        @time coeffs_nopar = Shearlab.SLsheardec2D(X_nopar, shearletSystem_nopar);  
        1.205391 seconds (7.12 k allocations: 1.354 GB, 26.23% gc time)
```


```
In [34]: # Compute the coefficients in parallel  
        X_par = AFArrray(convert(Array{Float32}, X_nopar));  
        @time coeffs_par = Shearlab.SLsheardec2D(X_par, shearletSystem_par);  
        0.273929 seconds (5.06 k allocations: 98.115 MB, 3.37% gc time)
```




```
In [35]: Shearlab.imageplot(real(Array(coeffs_nopar[:, :, 1])))
```




```
In [36]: Shearlab.imageplot(real(Array(coeffs_nopar[:, :, 5])))
```



```
In [37]: Shearlab.imageplot(real(Array(coeffs_nopar[:, :, 10])))
```



```
In [38]: Shearlab.imageplot(real(Array(coeffs_nopar[:, :, 16])))
```



Reconstruction of the image with the shearlet system and coefficients

```
In [41]: # Make the recovery
@time Xrec_nopar=Shearlab.SLshearrec2D(coeffs_nopar,shearletSystem_nopar);

0.969499 seconds (6.17 k allocations: 1.364 GB, 21.72% gc time)
```

```
In [42]: # Make the recovery in parallel
@time Xrec_par=Shearlab.SLshearrec2D(coeffs_par,shearletSystem_par);

0.010846 seconds (5.55 k allocations: 2.126 MB)
```

```
In [43]: # The recovery is very good
Shearlab.imageplot(Array(Xrec_nopar));
```



```
In [44]: # The recovery is very good
Shearlab.imageplot(Array(Xrec_par));
```



Benchmarks with matlab version


- 2D version.

Benchmark	Matlab(seconds)	Julia no gpu(seconds)	Julia gpu (seconds)	Improvement rate no gpu	Improvement rate gpu
Shearlet System 256x256	1.06	0.61	0.31	1.73	3.42
Decoding 256x256	0.18	0.15	0.043	1.2	4.19
Reconstruction 256x256	0.18	0.12	0.018	1.5	8.57
Shearlet System 512x512	5.15	3.07	1.8	1.22	2.08
Decoding 512x512	0.96	0.87	0.09	1.10	10.66
Reconstruction 512x512	0.84	0.52	0.021	1.62	14.00

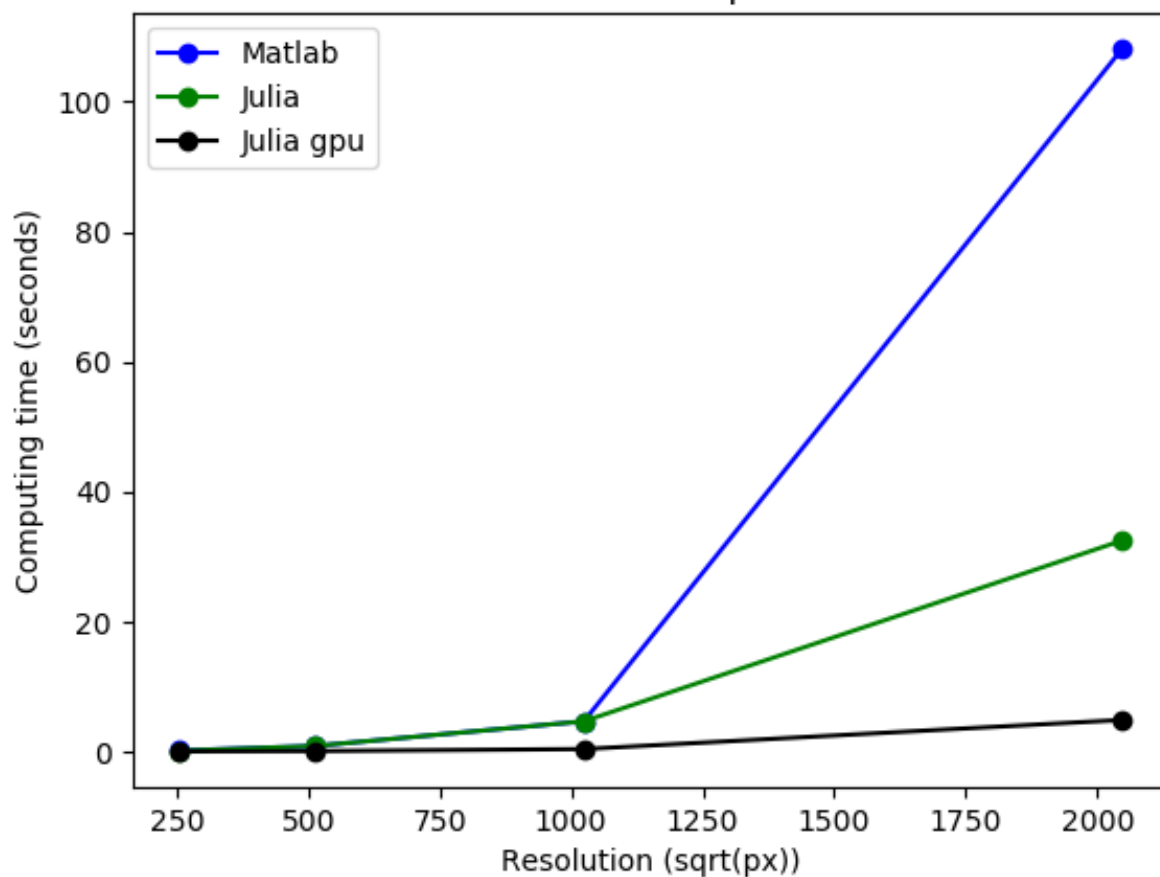
Benchmark	Matlab(seconds)	Julia no gpu(seconds)	Julia gpu (seconds)	Improvement rate no gpu	Improvement rate gpu
Shearlet System 1024x1024	35.84	15.65	2.44	2.29	14.68
Decoding 1024x1024	4.70	4.67	0.40	1.01	8.54
Reconstruction 1024x1024	4.72	4.48	0.037	1.05	127.56
Shearlet System 2048x2048	196.69	68.43	5.4	2.87	36.42
Decoding 2048x2048	108.19	32.50	5.88	3.33	18.39
Reconstruction 2048x2048	73.20	23.08	4.23	97.6	23.09

The benchmarks were made with 4 scales, in a Macbook pro with OSX 10.10.5, with 8GB memory, 2.7GHz Intel Core i5 processor and Graphic Card Intel Iris Graphics 6100 1536 MB.


- **Benchmarks plots 2D.**



Shearlet Decompostion



Shearlet Recovery



Applications

Image Denoising

Let $f \in \ell^2(\mathbb{Z}^2)$ and

$$f_{\text{noisy}}(i,j) = f(i,j) + e(i,j)$$

with noise $e(i,j) \sim \mathcal{N}(0, \sigma^2)$.

Then

$$f_{\text{denoised}} = S^* T_\delta S(f_{\text{noisy}})$$

where S is sparsifying transform (analysis operator of the shearlet system) and T_δ is the hard thresholding operator

$$(T_\delta x)(n) = \begin{cases} x(n) & \text{if } |x(n)| \geq \delta \\ 0 & \text{else.} \end{cases}$$

```
In [45]: # settings
        sigma = 30;
        scales = 4;
        thresholdingFactor = 3;
```

```
In [46]: # Give noise to data
        X_nopar = data_nopar;
        Xnoisy_nopar = X_nopar + sigma*randn(size(X_nopar));
```


```
In [48]: # Thresholding
        @time coeffs_nopar = coeffs_nopar.*(abs(coeffs_nopar).> thresholdingFactor*1
            size(X_nopar,1),size(X_nopar,2),length(shearletSystem_nopar.RMS))*si
```

0.534048 seconds (533.75 k allocations: 606.016 MB, 30.88% gc time)

```
In [49]: # Reconstruction
        Xrec_nopar = Shearlab.shearrec2D(coeffs_nopar, shearletSystem_nopar);
```

```
In [50]: # display
         clf;
         figure(figsize=(10,10));
         Shearlab.imageplot(Array(X_nopar), "Original", 1,2,1);
         Shearlab.imageplot(Array(Xnoisy_nopar), "Noised", 1,2,2);
```

Original



Noised



```
In [51]: elin = Shearlab.snr(Array(X_nopar),Array(Xrec_nopar));
         # display
         clf;
         figure(figsize=(10,10));
         Shearlab.imageplot(Array(X_nopar), "Original", 1,2,1);
         u = @sprintf("Denoised, SNR=%.3f", elin);
         Shearlab.imageplot(real(Array(Xrec_nopar)), u, 1,2,2);
```

Original



Denoised, SNR=27.405



Image inpainting

- Inpainting = optimization problem.

Let $x \in \ell^2(\mathbb{Z}^2)$ be a grayscale image partially occluded by a binary mask $\mathbf{M} \in \{0, 1\}^{\mathbb{Z} \times \mathbb{Z}}$, i.e.

$$y = \mathbf{M}x$$

One can recover x (inpaint) using an sparsifying transformation with the inverse problem

$$y^* = \min_{x \in \mathbb{R}^{N \times N}} \|S(x)\|_1 \text{ s.t. } y = \mathbf{M}x$$

By iterative thresholding algorithm

$$x_{n+1} = S^*(T_{\lambda_n}(S(x_n + \alpha_n(y - \mathbf{M}x_n))))$$

λ_n decreases with the iteration number linearly in $[\lambda_{min}, \lambda_{max}]$.

```
In [67]: function inpaint2D(imgMasked,mask,iterations,stopFactor,shearletsystem)
          coeffs = Shearlab.sheardec2D(imgMasked,shearletsystem);
          coeffsNormalized = zeros(size(coeffs))+im*zeros(size(coeffs));
          for i in 1:shearletsystem.nShearlets
              coeffsNormalized[:, :, i] = coeffs[:, :, i]./shearletsystem.RMS[i];
          end
          delta = maximum(abs(coeffsNormalized[:]));
          lambda=(stopFactor)^(1/(iterations-1));
          imgInpainted = zeros(size(imgMasked));
          #iterative thresholding
          for it = 1:iterations
              res = mask.*(imgMasked-imgInpainted);
              coeffs = Shearlab.sheardec2D(imgInpainted+res,shearletsystem);
              coeffsNormalized = zeros(size(coeffs))+im*zeros(size(coeffs));
              for i in 1:shearletsystem.nShearlets
                  coeffsNormalized[:, :, i] = coeffs[:, :, i]./shearletsystem.RMS[i];
              end
              coeffs = coeffs.*(abs(coeffsNormalized).>delta);
              imgInpainted = Shearlab.shearrec2D(coeffs,shearletsystem);
              delta=delta*lambda;
          end
          imgInpainted
      end
```

WARNING: Method definition inpaint2D(Any, Any, Any, Any, Any) in module Main at In[66]:2 overwritten at In[67]:2.

Out[67]: inpaint2D (generic function with 1 method)

```
In [68]: # Rename images
          img_nopar = data_nopar
          img_par = data_par;
```

```
In [69]: # Import two masks
          name = "../.../ShearLab.jl/data_samples/data_inpainting_2d/mask_rand.png";
          mask_rand = Shearlab.load_image(name, n);
          mask_rand = mask_rand[:, :, 1];
          name = "../.../ShearLab.jl/data_samples/data_inpainting_2d/mask_squares.pr
          mask_squares = Shearlab.load_image(name, n);
          mask_squares = mask_squares[:, :, 1];
```

```
In [70]: # Setting of data
imgMasked_rand_nopar = img_nopar.*mask_rand;
imgMasked_rand_par = AFArrray(convert(Array{Float32},img_par.*mask_rand));
imgMasked_squares_nopar = img_nopar.*mask_squares;
imgMasked_squares_par = AFArrray(convert(Array{Float32},img_par.*mask_squares);
stopFactor = 0.005; # The highest coefficient times stopFactor
sizeX = size(imgMasked_rand_par,1);
sizeY = size(imgMasked_rand_par,2);
nScales = 4;
shearLevels = [1, 1, 2, 2];
```

```
In [71]: tic()
imginpainted_rand50_nopar = inpaint2D(imgMasked_rand_nopar,mask_rand,50,stopFactor);
toc()
```

elapsed time: 158.482325353 seconds


Out[71]: 158.482325353

```
In [72]: clf;
figure(figsize=(10,10));
Shearlab.imageplot(img_nopar, "Original", 1,2,1);
Shearlab.imageplot(imgMasked_rand_nopar, "Masked random", 1,2,2);
```

Original



Masked random



```
In [73]: elin = Shearlab.snr(img_nopar,imginpainted_rand50_nopar);  
# display  
clf;  
figure(figsize=(10,10));  
Shearlab.imageplot(img_nopar, "Original", 1,2,1);  
u = @sprintf("Inpainted random 50 iterations, SNR=%.3f", elin);  
Shearlab.imageplot(real(imginpainted_rand50_nopar), u, 1,2,2);
```

Original



Inpainted random 50 iterations, SNR=24.423



```
In [74]: tic()  
imginpainted_squares50_nopar = inpaint2D(imgMasked_squares_nopar,mask_square  
toc()
```

elapsed time: 153.952855039 seconds


Out[74]: 153.952855039

```
In [75]: clf;
figure(figsize=(10,10));
Shearlab.imageplot(img_nopar, "Original", 1,2,1);
Shearlab.imageplot(imgMasked_squares_nopar, "Masked squares", 1,2,2);
```

Original



Masked squares



```
In [76]: elin = Shearlab.snr(img_nopar,imginpainted_squares50_nopar);
# display
clf;
figure(figsize=(10,10));
Shearlab.imageplot(img_nopar, "Original", 1,2,1);
u = @sprintf("Inpainted squares 50 iterations, SNR=%.3f", elin);
Shearlab.imageplot(real(imginpainted_squares50_nopar), u, 1,2,2);
```

Original



Inpainted squares 50 iterations, SNR=28.182



```
In [ ]:
```

